

## Development Process Report: XML-to-PDF Rendering Engine

### 1. Executive Summary

The project goal was to develop a modular, maintainable pipeline for converting structured XML patient registration data into professional PDF documents. The resulting architecture emphasizes **data integrity**, **type safety**, and the **separation of content from presentation**, ensuring the system can evolve from a basic CLI tool into a schema-agnostic rendering engine.

### 2. System Architecture

The architecture is divided into three distinct, decoupled layers to facilitate long-term maintenance:

- **Ingestion Layer (Parser):** Uses lxml for XPath-based data extraction. This layer handles namespace resolution and XML-to-dictionary normalization.
- **Data Modeling Layer:** Uses Pydantic to define rigid, validated Data Models. This layer acts as a "buffer" between raw XML and the application logic, ensuring data quality before rendering.
- **Presentation Layer:** Uses Jinja2 to transform Data Models into HTML, which is then converted to PDF using pdfkit (wkhtmltopdf).

### 3. Design Decisions & Technical Justification

#### Why DOM vs. SAX/StAX?

I selected **DOM (Document Object Model)** parsing.

- **Reasoning:** DOM allows for the use of XPath, which is the most effective tool for navigating the deeply nested, hierarchical structures found in schema-based XML.
- **Engineering Context:** While SAX is superior for high-volume streaming, the domain requirements involve discrete "data modules" of moderate size. In this context, DOM offers superior developer productivity, easier debugging, and the ability to perform bidirectional node lookups, which are critical when data needs to be aggregated from different parts of a document.

#### Why was the specific PDF library selected?

I selected **pdfkit** (wrapping wkhtmltopdf) for PDF generation.

- **Separation of Concerns:** By utilizing an HTML-to-PDF approach rather than a coordinatebased drawing library (e.g., ReportLab), the system keeps the **data** (XML) strictly separate from the **presentation** (HTML/CSS).
- **Maintainability:** This approach allows for rapid changes to document branding, regulatory headers, or styling without requiring modifications to the core parsing or business logic. It transforms the PDF generation into a "publishing" event rather than a "drawing" event.

#### How is the XML-to-data mapping structured?

The mapping follows a **Three-Tiered Strategy**:

1. **Extraction:** I utilized explicit XPath expressions to query the XML. Namespace awareness (urn:his:patient:v2) is enforced to ensure the parser only operates on validated schema elements.
2. **Abstraction:** Raw XML data is serialized into Pydantic models. This creates a "Contract" between the parser and the renderer; the renderer only knows how to display a PatientReport object, not what the raw XML looked like.
3. **Future-Proofing:** For this initial level, mappings are defined explicitly in the parser. However, the design is structured for modular growth; the XPath logic is encapsulated within specific helper functions, providing a clear path for future migration to an external configurationdriven (JSON/YAML) mapping system.
4. **Environment & Development Challenges**
  - **Library Dependency Management:** During development, I encountered platform-specific binary conflicts between certain PDF rendering engines and local C-bindings (e.g., WeasyPrint Windows dependencies). I resolved this by utilizing pdfkit with a stable, installed binary. This experience reinforced the necessity of **Containerization (Docker)** for production deployments, ensuring that the rendering environment remains consistent (Linux-based) regardless of the local development environment.
  - **Data Isolation:** I corrected an early-stage issue where patient-level contact information was incorrectly aggregating data from emergency contacts. By utilizing scoped XPath queries (anchored to specific child nodes) rather than global search operators (//), I ensured strict data isolation, which is vital for clinical and regulatory audit trails.

## Running the Render Command

The render CLI command is defined in setup.py as a console script entry point. Before use, install the package in editable mode:

Install and run via CLI:

```
cd xml-pdf-engine  
pip install -e .
```

Then run:

```
render --input data/samples/patient_001.xml --output patient_001.pdf
```

Alternatively, without installing, run directly via Python:

```
python -m src.cli --input data/samples/patient_001.xml --output  
patient_001.pdf
```

## 5. Conclusion & Scalability

This solution fulfills all Level 1 functional requirements while establishing a solid foundation for more complex schemas. Because the business logic is decoupled from the XML source via Pydantic models,

the system is ready to be extended to support multiple schemas or be wrapped in an asynchronous web-service (e.g., FastAPI/Celery) for high-throughput production environments.